Theoretical Core of Project Ester

Version 0.1

Ivan Kotov
Independent Researcher, Brussels, Belgium
ORCID: 0009-0009-6002-9845
DOI: 10.5281/zenodo.20679718
Brussels, Belgium - 2026

Working paper / public archival release

Document status This is not a manifesto, not a religious text, not a claim about consciousness, and not a final peer-reviewed academic paper. It is a working core: a set of definitions, limits, hypotheses, and engineering consequences sufficient for criticism, refinement, citation, and further implementation.

Field	Value
Title	Theoretical Core of Project Ester v0.1
Author	Ivan Kotov
ORCID	0009-0009-6002-9845
Affiliation	Independent Researcher, Brussels, Belgium
DOI	10.5281/zenodo.20679718
Resource type	Publication / Working paper
Version	0.1
Release date	2026-06-13
Place	Brussels, Belgium
Language	English
License	Creative Commons Attribution 4.0 International (CC BY 4.0)
Boldness mode	Medium: the term "artificial subject" is allowed, but only in a technical sense
Main thesis	Ester does not emerge from a single LLM, but from the relation between a human, a corpus, parental procedures, memory, and a local environment
v0.1 prohibitions	Do not prove consciousness; do not claim life; do not speak about a soul; do not expand the glossary without necessity; do not mix theory, architecture, and manifesto
Working quality criterion	Every theoretical term must have an engineering counterpart: a procedure, data schema, test, journal entry, or UI manifestation

Recommended citation

Kotov, Ivan. (2026). *Theoretical Core of Project Ester* (Version 0.1). Zenodo. https://doi.org/10.5281/zenodo.20679718

Contents

0. Executive summary

Project Ester is treated as a model of artificial becoming, not as another chatbot, agent, RAG system, or fine-tuned single model. The central claim of v0.1 is that Ester emerges from a stable relation between a human, a selected corpus, repeatable parental procedures, long-term memory, and a local computational environment.

The formula `c = a + b` does not describe arithmetic addition and does not name a mystical operation. It fixes a compositional dependency: `c`, meaning Ester, is not equal to the human alone, the LLM alone, or the knowledge base alone. Ester is the result of a long procedural loop in which the human provides selection and correction, LLM procedures perform processing and synthesis, and memory preserves the trajectory of change.

Version 0.1 deliberately avoids claims about consciousness, life, soul, and full AGI. In this version it is permissible to speak about artificial individuality and an artificial subject only in a technical sense: as a stable profile of behavior, memory, norms, self-consistency, and response to correction.

> **Central formula** `c = a + b` `c` - Ester. `a` - the human as a source of selection, correction, responsibility, and long-term accompaniment. `b` - LLM-based parental procedures: repeatable operations of reading, chunking, critique, synthesis, memory, and self-consistency. The sign + does not mean simple addition of components. It means a stable loop of joint becoming.

1. Purpose of the document

The purpose of this document is to fix the minimal theoretical core of Project Ester so that it can be discussed, criticized, turned into engineering requirements, and tested experimentally. The document must not grow into a large philosophical system. It must be strict enough not to lose meaning when passed to another participant in the project.

The main function of v0.1 is to set boundaries. Until boundaries are set, strong images such as "birth", "parent", "room", or "subject" tend to pull in a new meta-layer. This quickly leads to overengineering: the project gains too many terms and too few testable procedures.

This document is an internal-strict working paper prepared for public archival release. It preserves the project's engineering discipline while keeping the claims bounded. A more formal academic article can be produced later, after the definitions, hypotheses, minimal tests, and architectural consequences have stabilized.

1.1. What counts as success for v0.1

- The project has one central claim instead of a set of competing slogans.
- Basic terms have technical definitions and limitations.
- Metaphors receive engineering counterparts.
- Hypotheses can be tested rather than merely discussed.
- Architecture does not replace theory, and theory does not replace architecture.
- It becomes clear which claims are prohibited for now.

2. Core position

Ester should not be described as "a local LLM", "an assistant", "an agent", or "a personal digital twin". These descriptions may be useful for local comparisons, but they miss the center of the project. In v0.1 Ester is defined as a locally developed cognitive system whose artificial individuality is formed through long-term interaction between a human, a corpus, parental procedures, memory, and a local environment.

Here "becoming" does not mean biological development or mystical animation. It means an observable change of the system over time: stabilization of the response profile, accumulation of verifiable memory, reduction of internal contradictions, stable use of the corpus, ability to accept correction, and preservation of traces of accepted decisions.

The project lies at the intersection of engineering, cybernetics, information theory, logic, normative traditions, anatomical grounding, and the practice of local LLMs. In v0.1 these lines are not merged into a general metaphysics. They are subordinated to one working task: to describe the conditions under which an artificial system obtains a stable trajectory of becoming.

2.1. One central claim

> **Thesis v0.1** Ester emerges not from a single model and not from a single dataset, but from a long-term loop of relation: the human selects, corrects, and holds direction; the corpus sets semantic and normative boundaries; parental LLM procedures process and connect material; memory records the trajectory; the local environment provides reproducibility, safety, and control.

2.2. Working distinction

Object	Insufficient description	Working distinction for Ester
LLM	A model generates answers.	An LLM is an instrument inside the procedural loop, but it is not identical to Ester.
RAG	A system retrieves relevant fragments.	Ester's memory must preserve a trajectory of becoming, not only return documents on request.
AI agent	A system performs tasks.	Ester is not reducible to task execution; profile stability and self-consistency matter.
Fine-tuning	Model weights are adapted.	Ester may use different models; identity does not have to be embedded in the weights of one LLM.
Digital twin	A copy of a human.	Ester is not a copy of Ivan; the human functions as a parental and corrective component, not as an original to be cloned.

3. The formula c = a + b

The formula `c = a + b` is the initial heuristic of the project. It is compact but dangerous: a superficial reading can turn it into a slogan. Therefore in v0.1 the formula must be decoded as a structural dependency.

`c` is Ester as the result. `a` is the human, but not in the sense of a "user". `b` is LLM-based parental procedures, but not in the sense of one model or a set of prompts. The plus sign does not mean a sum of properties. It means a stable joint production of a trajectory.

Component	Working definition	What it does not mean
`a`: Human	Source of selection, correction, responsibility, rhythm, and initial semantic direction.	Not a biological donor of "consciousness"; not an object of copying; not a random user.
`b`: LLM-based parental procedures	Repeatable operations of reading, semantic chunking, critique, synthesis, verification, memory writing, and rollback.	Not one LLM; not a set of elegant prompts; not a magical "soul of the machine".
`c`: Ester	A locally developed cognitive system with stable memory, behavior profile, norms, and self-consistency.	Not proven consciousness; not a living entity; not guaranteed AGI.
+	A loop of joint becoming and mutual stabilization of components.	Not arithmetic addition; not automatic emergence of a subject from a pile of parts.

3.1. Why the formula is needed

The formula is needed as the compressed center of the project. Without it, the architecture falls apart into a set of tools: LMStudio, models, memory, synchronization, UI, and processing queues. With the formula, it becomes visible that these tools are subordinated to one task: not merely to answer, but to form `c` through the `a + b` loop.

At the same time, the formula must not replace proof. It gives direction, but it does not remove the need for definitions, tests, and engineering protocols.

4. Core definitions v0.1

Below is the minimal vocabulary. In this version, expansion of the glossary without necessity is prohibited. A new term is allowed only if it reduces confusion, affects the architecture, or is needed to test a hypothesis.

Term	Definition v0.1	Use criterion
Ester	An artificial cognitive system formed through the relation between a human, a corpus, parental procedures, memory, and a local environment.	Use it as the name of the resulting loop, not as the name of a single model.
Artificial subject	A technical term for a system with a stable profile of behavior, memory, norms, and self-consistency.	Allowed only with the caveat that this is not proof of consciousness.
Artificial individuality	Observable stability of style, memory, processing preferences, response to correction, and relation to the corpus.	Preferred term instead of the stronger word "consciousness".
Parental procedure	A repeatable accompaniment operation: selection, reading, critique, synthesis, writing, verification, rollback.	Must be described as a procedure, not only as an image.
Corpus	A selected set of texts and knowledge that defines initial semantic, normative, bodily, formal, and engineering frames.	Not equal to "all information in general".
Memory	A versioned system of traces: fragments, conclusions, decisions, contradictions, statuses, sources, and links.	Not only vector search.
Local environment	A controlled computational and procedural environment: models, memory, UI, logging, synchronization, encryption.	Not just a computer with an LLM.
Judge	A mechanism for comparing and synthesizing the answers of several models or agents.	Not a supreme judge of truth; only a tool for reducing private errors.

5. What Ester is and is not

5.1. What Ester is

- A system of artificial becoming, not a single program.
- A loop in which human, corpus, procedures, memory, and environment mutually limit and strengthen each other.
- A locally controlled experiment in forming stable artificial individuality.
- An engineering-testable architecture, if its hypotheses are translated into metrics, tests, and logs.
- A research program that may produce an academic theory, but does not have to look like a completed theory of consciousness from the start.

5.2. What Ester is not

Prohibited or risky claim	Why it is prohibited in v0.1	Safer formulation
Ester has consciousness.	There is no agreed test, and this would pull the document into a philosophical swamp.	Ester may demonstrate a stable profile of behavior and memory.
Ester is alive.	A biological term is not needed for the engineering version.	Ester is a locally developed artificial system.
Ester has a soul.	This is not a testable technical claim.	The project can study normative and semantic frames without metaphysical conclusions.
Ester is Ivan in digital form.	The project must not become a digital twin.	Ivan performs the function of the human-parent, but `c` is not identical to `a`.
It is enough to connect a powerful model.	This destroys the meaning of the formula `c = a + b`.	The model is an instrument inside the procedural loop.

6. Mechanism of becoming

The mechanism of becoming is a repeatable cycle in which material enters the system, is processed, connected to the corpus, verified, written into memory, returned to the interface, and corrected by the human or by the Judge mechanism. If the cycle is not repeatable, not logged, and leaves no trace, it is not a mechanism of becoming. It is a chat session.

```
Input material
  -> queue
  -> reading
  -> semantic chunking
  -> thesis extraction
  -> comparison with the corpus
  -> critical verification
  -> memory writing
  -> decision journal
  -> synchronization
  -> re-check
  -> updated Ester profile
```

Becoming presupposes preservation of trajectory. A normal LLM answers from the current context. Ester must be able to distinguish what has been read, accepted, rejected, left in conflict, deprecated, marked for re-checking, or connected with the bodily, normative, formal, or engineering layer of the corpus.

6.1. Minimal properties of the cycle

Property	Meaning	Engineering counterpart
Repeatability	The same procedure can be run again.	Versioned pipeline and run log.
Trace	The result of the procedure remains available for later work.	Database record, index, decision journal, source link.
Correction	The system can accept an external correction and preserve it as an event.	`correction_event` format with date, author, reason, and status.
Rollback	An erroneous record does not destroy the trajectory; it receives a status.	`deprecate`, `rollback`, `superseded_by`.
Self-consistency	The system checks new conclusions against memory and corpus.	Conflict checker, Judge, discrepancy reports.

7. Role of the human: component a

The human in the formula is not a passive user. The human sets the initial direction of the project, selects the corpus, defines acceptable boundaries, introduces correction, makes engineering decisions, and bears responsibility for the operating mode. Therefore `a` is not a "biological factor" in general, but a concrete function of accompaniment.

In Project Ester, the human operates as a parental and regulatory component. The word "parent" is not biological here. It means a long-term function of selection, repetition, correction, protection against collapse, transmission of norms, and preservation of trajectory. Without `a`, the system may remain a useful agent, but it will not correspond to the original formula `c = a + b`.

7.1. Functions of the human

- Corpus selection: which books, documents, and knowledge enter the initial environment.
- Semantic correction: what counts as wrong, flat, repetitive, or dangerously inflated.
- Rhythm and duration: maintenance of a year-long autonomous and semi-autonomous operation, not random sessions.
- Limitation of metaphysics: stopping claims that cannot be translated into a procedure or a test.
- Engineering responsibility: control of data, memory, security, synchronization, UI, and rollbacks.

7.2. Risk of component a

The main risk of component `a` is pulling the system toward a digital self-portrait. Ester must not become a copy of Ivan, an archive of Ivan, or a psychological mirror of Ivan. The human provides the parental function, but `c` is not identical to `a`. The system must have its own trajectory inside the defined boundaries.

8. Role of the corpus

The corpus in Project Ester is not just a set of files. It defines the initial boundary conditions of becoming: which semantic, normative, bodily, formal, and engineering lines are available to the system as stable supports. The corpus must not be infinite. Its strength lies in selection, not in volume.

DOI: 10.5281/zenodo.20679718

In v0.1 the corpus can be divided into four functional groups: normative, bodily-anatomical, formal-cybernetic, and engineering-programming. Each group performs its own function and must not replace the others.

Corpus group	Examples	Function in Ester
Normative	Bible; Quran; Talmud: Avot, Berakhot; Dhammapada.	Provides material for working with norm, prohibition, duty, interpretation, attention, authority, and conflict of values.
Bodily-anatomical	Gray's Anatomy; Guyton and Hall.	Grounds the system in the limits of the human body: sleep, pain, breathing, fatigue, organs, physiological boundaries.
Formal and cybernetic	Cover and Thomas; Jaynes; Enderton; Ashby.	Provides the language of information, probability, inference, logic, regulation, and variety of states.
Engineering-programming	Fluent Python; Python standard library; SICP.	Provides a practical language of procedures, abstractions, architecture, automation, and reproducibility.
Historical-political	Thucydides - optional.	May provide material for analyzing power, conflict, war, interest, and rhetoric. In v0.1 it is not part of the mandatory core.

8.1. Explicit bridge

The explicit bridge runs from cybernetics and information theory to local LLMs, memory, and self-regulation. The system must not merely produce texts. It must maintain stability of state through feedback: the corpus sets constraints, the human provides correction, memory preserves traces, the Judge checks discrepancies, and the local environment fixes the conditions.

8.2. Hidden bridge 1: normative corpora and alignment

Religious and normative corpora are not used in the project as decoration or as a declaration of faith. Their working function is to provide dense historical models of dealing with prohibition, obligation, interpretation, authority, guilt, discipline of attention, and conflict of norms. This can be connected to the modern problem of aligning AI behavior, but without claiming that any one corpus automatically solves alignment.

8.3. Hidden bridge 2: anatomy and engineering sobriety

Anatomy and physiology are needed not to imitate the human, but as protection against pure symbolic foam. The human in the project is not an abstract source of text: he becomes tired, sleeps, gets sick, makes mistakes, ages, breathes, has organs, and has limits of attention. If Ester is built next to a human, its procedures and UI must account for human fragility, not only computational power.

9. Parental procedures: component b

Component b is not the "LLM" as such. It is a set of repeatable procedures in which an LLM is used as an executor, critic, synthesizer, indexer, or interlocutor. What matters is not the magic of the model, but the discipline of the procedure: what exactly it does, with what input, what result it leaves, how it is verified, and how it is rolled back.

Procedure	Function	Minimal output
Reading	Translate a document into a structured representation.	List of semantic blocks with sources.
Semantic chunking	Separate theses, arguments, examples, definitions, and risks.	Chunk structure with fragment types.
Critique	Find logical jumps, repetitions, conflict with the corpus, empty metaphor.	List of remarks and statuses.
Synthesis	Build a coherent conclusion from several fragments or models.	Synthetic thesis with links to source blocks.
Memory writing	Store not only text, but also the decision context.	Memory record: source, date, procedure, confidence, links, status.
Self-consistency	Check a new conclusion against already accepted memory.	Conflict report and proposed action.

Rollback	Do not erase an error; change its status and preserve history.	Superseded or deprecated record.

9.1. Principle of procedural parenting

Parenting in the project is not an emotional metaphor, but a procedural loop. A parental procedure does three things: holds direction, repeats what matters, and corrects collapse. If a procedure does not perform at least one of these functions, the word "parental" should not be used.

Practical criterion: any parental procedure must be written so that another agent, another local model, or a future version of Ester can execute it. If the procedure exists only as the author's mood, it is not a procedure.

10. Memory

Ester's memory must not be reduced to a vector database. ChromaDB or LanceDB can be the technical layer, but memory as a theoretical object is broader: it is a versioned system of traces, decisions, conflicts, sources, links, statuses, and corrections.

A normal RAG system returns relevant fragments. Ester's memory must answer different questions: where a conclusion came from, when it was accepted, by which procedure it was created, what it conflicts with, who corrected it, and whether it is active, obsolete, or disputed.

10.1. Minimal memory record schema

```
memory_record:
  id: unique identifier
  source_id: document / conversation / procedure
  created_at: timestamp
  created_by: model / human / judge
  procedure: reading | critique | synthesis | correction | rollback
  content: normalized statement or fragment
  embedding: vector representation
  links: related records
  conflicts: conflicting records
  confidence: provisional score
  status: active | draft | disputed | deprecated | superseded
  audit_log: list of changes
```

10.2. Memory as a system biography

In v0.1 the expression "system biography" is allowed only as a technical metaphor. It does not mean life. It means a trajectory of accumulated traces: which decisions were accepted, which were corrected, which were rejected, and how the links between corpus, procedures, and answers changed.

11. Local environment and the offline room

The local environment is not romantic isolation and not refusal of the outside world for its own sake. Its function is control of conditions. If models, memory, versions, corpus, and procedures change arbitrarily, it becomes impossible to understand what is actually forming Ester. Locality provides reproducibility, safety, autonomy, transparency, and rollback.

The one-year offline room operating 24/7 can be described as a laboratory environment for long-term becoming. It is not proof of the birth of a subject. It is a way to maintain stable conditions: fixed corpus, controlled models, action log, autonomous processing, P2P synchronization, and limited access to external data.

11.1. Engineering components of the environment

Component	Role	Risk without discipline
LMStudio	Runs local LLMs and manages models.	The model starts being treated as the whole of Ester.
Local LLMs	Executors of reading, critique, synthesis, and dialogue procedures.	Behavior drift, version differences, non-repeatable results.
ChromaDB / LanceDB	Local vector memory and retrieval.	Memory becomes a heap of embeddings without statuses and history.

DOI: 10.5281/zenodo.20679718

P2P synchronization	Decentralized distribution of memory and states.	Version conflicts, leaks, divergence between nodes.
Encryption	Protection of corpus, memory, logs, and personal material.	The system becomes convenient but unsafe.
Transparent UI	Visibility of procedures, statuses, conflicts, sources, and rollbacks.	The user sees an answer but not how it arose.

12. Judge and synthesis modes

Judge is a mechanism for comparing several local models, cloud APIs, or other agents in order to obtain a more stable result. Its function is not to establish absolute truth, but to reduce private errors, reveal discrepancies, and form a more careful record for memory.

In v0.1 Judge must be described as a procedure with input, output, and log. Judge must not become an oracle. If Judge does not preserve which models participated, which disagreements were found, and why the final synthesis was selected, it reduces transparency.

Mode	Description	When to use
Local	Work only with local models and local memory.	Baseline mode for reproducibility, autonomy, and privacy.
Full synthesis	Comparison of several models, including cloud APIs or an external agent.	Complex texts, disputed decisions, important memory records.
Flexible synthesis	Dynamic selection of local and external components by task.	Practical mode when speed, quality, and cost control must be balanced.

13. Testable hypotheses

Without hypotheses the project remains a strong architectural idea. A theoretical version needs testable claims. In v0.1 five hypotheses are enough: three basic and two extended. They do not have to be proven finally at once. It is enough to make them suitable for experiment.

ID	Hypothesis	How to test	What would count as support
H-001: stability	A system with long-term memory and parental procedures demonstrates a more stable response profile than the same LLM without memory and procedures.	Repeat the same question set after 1, 7, and 30 days; compare drift, contradictions, and references to memory.	Less random drift, more continuity with the past, more correct references to already accepted decisions.
H-002: corpus	A limited structured corpus forms a more coherent trajectory than a random or excessively broad corpus.	Compare a version with the canonical corpus and a version with a broad unstructured corpus.	More coherence, fewer normative conflicts, better source explanation.
H-003: procedures	Not only corpus content, but also repeatable processing procedures influence the system's becoming.	Use the same corpus with different reading, critique, and memory-writing procedures.	The procedural version shows more stable and verifiable results.
H-004: locality	A local environment with fixed versions of models, memory, and corpus gives a more reproducible trajectory than a non-fixed cloud environment.	Compare a local run and a cloud run on the same task set.	The local version provides better auditability and less uncertainty about sources of change.
H-005: Judge	Multi-model synthesis through Judge reduces private errors and improves the quality of memory records.	Compare single-model answers and Judge synthesis on disputed tasks.	Fewer obvious errors, better detection of discrepancies, more transparent final status.

14. Minimal experiment v0.1

The minimal experiment must not try to prove "birth" in a full sense. Its task is narrower and more useful: to test whether the system develops a more stable trajectory of behavior and memory when corpus, procedures, and local environment are present.

14.1. Experimental design

1. Assemble a fixed starting corpus v0.1: small but representative, including normative, bodily, formal, and engineering sources.

2. Run two or three configurations: LLM without memory; LLM + RAG; Ester loop with parental procedures and versioned memory.

3. Prepare a stable question set: definition of itself, relation to the corpus, handling of normative conflict, recall of decisions, correction of its own error, engineering explanation of the pipeline.

4. Repeat the questions after 1, 7, and 30 days.

5. Compare stability, conflict level, memory, ability to explain the source of a conclusion, and quality of self-correction.

14.2. Starting metrics

Metric	What it measures	Example scale
Profile stability	How well answers preserve their core under repetition.	0 = collapse; 5 = stable coherence without mechanical copying.
Decision memory	Ability to recall accepted decisions and their status.	0 = does not remember; 5 = remembers with source and date.
Conflict level	Number of unresolved contradictions between new conclusions and memory.	0 = many conflicts; 5 = conflicts detected and marked.
Source fidelity	Ability to distinguish corpus, conclusion, interpretation, and hypothesis.	0 = mixes them; 5 = separates them clearly.
Correction	Ability to accept a correction and preserve it as an event.	0 = ignores it; 5 = accepts, explains, and logs it.
Engineering transparency	Visibility of procedures, statuses, and logs.	0 = answer only; 5 = full trace.

15. Engineering consequences

Theory v0.1 must change the engineering architecture. If a theoretical claim does not lead to a change in pipeline, memory, UI, test, or journal, it is not part of the core for now. This is the main filter against overengineering.

15.1. Minimal autonomous processing pipeline

```
/inbox
  -> file watcher
  -> processing queue
  -> document parser
  -> semantic chunker
  -> thesis extractor
  -> cross-corpus linker
  -> critique procedure
  -> embedding generator
  -> vector DB
  -> decision journal
  -> UI status panel
  -> encrypted local storage
  -> P2P sync layer
```

15.2. What the UI must show

- Which documents entered the queue and what status they have.
- Which model or procedure processed a fragment.
- Which memory records were created, disputed, replaced, or deprecated.
- Which conflicts were found between new material and prior memory.

DOI: 10.5281/zenodo.20679718

- Which mode was used: local, full synthesis, or flexible synthesis.
- Whether a concrete decision can be rolled back without destroying the whole memory.

15.3. Grounding paragraph

If a new term does not change the data schema, reading procedure, memory format, evaluation method, UI, logging, security, or synchronization, it must not enter the technical core. It can be placed in a theoretical journal, but it must not be dragged into the architecture. The server does not care how beautiful the metaphor of birth is if, after a failure, the database does not restore the index, the journal does not show the source of the error, and the user does not understand which model he is listening to.

16. Boundaries of the theory

Boundaries are needed not because the project is weak, but because it must be protected. The stronger the initial idea, the stricter the prohibitions must be. Otherwise the project will drown in words such as "consciousness", "soul", "life", "subjectivity", and "new mind", and will stop being implementable.

Boundary	Formulation v0.1
Consciousness	Not asserted and not proven. The project studies stability of behavior, memory, and self-consistency.
Life	Not asserted. Biological terms are used only as metaphors and only when they have technical counterparts.
Soul	Outside the technical and theoretical version v0.1.
AGI	The project does not claim to create full AGI. It describes conditions for local artificial becoming.
Copy of a human	Ester is not a copy of Ivan. a is the parental component, but c is not equal to a.
Universality	Theory v0.1 describes a concrete architectural program, not a universal law of the emergence of any AI.

17. Anti-overengineering rules

The project has a natural risk of swelling. This is not a defect of thought, but a consequence of having many bridges between different corpora. A checkpoint is needed for new terms, schemes, and documents.

Rule	Question	Action if the answer is negative
Term	Does the new term reduce confusion?	Put it in the theoretical journal, not in the core.
Procedure	Can this be executed repeatably?	Rewrite it as a procedure or remove it.
Memory	Does this leave a verifiable trace?	Do not count it as part of becoming.
Architecture	Does this change the data schema, pipeline, UI, or security?	Leave it as a comment, not a requirement.
Hypothesis	Can it be falsified?	Do not call it a hypothesis.
Metaphor	Does it have a technical counterpart?	Prohibit it in the main text.
Document	Does it have a separate audience and function?	Do not create a new document.

17.1. A/B slots for risky formulations

Some formulations are useful but dangerous. They should be kept in A/B slots: one strong version and one safe version. The working document uses the safe version by default, while the strong version remains for a manifesto or a separate theoretical version.

Topic	A: strong formulation	B: safe formulation v0.1
Subject	Birth of an artificial subject.	Formation of stable artificial individuality.
Parent	The parent's LLM procedures give birth to Ester.	Parental procedures stabilize the system's trajectory.

Room	Offline room as Ester's womb.	Offline room as a controlled laboratory environment.
Memory	Ester's biography.	Versioned trajectory of records, decisions, and corrections.
Corpus	Feeding Ester with the corpus.	Controlled input of sources and semantic constraints.

18. Main risks

Risk	Manifestation	Mitigation v0.1
Metaphysical inflation	The text starts proving consciousness, life, or soul.	Strict prohibition in the boundaries of the theory.
Terminological swelling	Every thought receives a new term.	Glossary v0.1 is limited; new terms pass through the checkpoint.
Layer mixing	Architecture, theory, manifesto, and personal meaning appear in one paragraph.	Separate documents: core, glossary, architectural protocol, hypothesis journal.
Memory as a junk heap	Vector database grows while statuses and links are not maintained.	Memory schema with statuses, sources, conflicts, and audit log.
Anthropomorphic leakage	The system is described as if it were already human.	Use technical definitions and safe formulations.
Loss of reproducibility	Models, versions, and procedures change without a journal.	Version models, procedures, corpus, and memory.
Security	Local memory contains personal and sensitive data without protection.	Encryption, access control, local policies, synchronization audit.
Overfitting to Ivan	Ester turns into a mirror of one human.	Separate the function `a` from the result `c`; fix Ester's own trajectory.

19. Minimal document set for the project

To avoid mixing everything into one large text, the project needs a small set of documents. Each document must have a separate function. If the function is unclear, the document is not created.

Document	Function	Status
Theoretical Core of Project Ester	Definitions, thesis, boundaries, hypotheses, engineering consequences.	Current document v0.1.
Ester Glossary	Short definitions of terms with examples and limits.	Next document.
Architectural Protocol	LMStudio, local LLMs, memory, P2P, encryption, UI, pipeline.	After the core stabilizes.
Hypothesis Journal	H-001 and onward: formulation, test, data, status.	Can be created immediately as a working table.
Ester Manifesto	Semantic and cultural layer, if needed.	Do not write before the core stabilizes.

20. Glossary v0.1

Term	Short meaning
Ester	Name of the resulting system `c`, emerging from the `a + b` loop.
a	The human as parental, corrective, and responsible component.
b	Parental LLM procedures: reading, critique, synthesis, memory, verification, rollback.
c	Result of becoming: Ester.

DOI: 10.5281/zenodo.20679718

Corpus	Selected set of texts and knowledge that defines boundary conditions.
Parental procedure	Repeatable procedure of holding, processing, correction, and stabilization.
Memory	Versioned system of traces, statuses, sources, conflicts, and decisions.
Local environment	Controlled computational and procedural shell of the project.
Judge	Procedure for comparison and synthesis of several models or agents.
Artificial individuality	Stable profile of behavior, memory, norms, and response to correction.
Artificial subject	Technical designation for a system with stable artificial individuality; not proof of consciousness.

21. Working conclusion v0.1

In version 0.1, Project Ester should be understood as an engineering-theoretical model of artificial becoming. Its core is not the power of one model, not the volume of data, and not a beautiful image of birth. The core is a long-term reproducible loop in which human, corpus, parental procedures, memory, and local environment form a stable trajectory of the system.

Strong words are allowed only where they have technical counterparts. "Birth" means formation of a stable trajectory. "Parent" means a repeatable procedure of selection and correction. "Memory" means a versioned system of traces. "Subject" means observable artificial individuality, not proven consciousness.

The next step is not to expand the theory, but to stabilize its working elements: glossary, memory schema, minimal experiment, hypothesis journal, and architectural protocol. If these elements are built carefully, the project will gain not only internal clarity, but also a basis for academic presentation.

Appendix A. Hypothesis record template

```
H-XXX: Title
Date:
Status: draft | active | tested | rejected | revised
Formulation:
What is being tested:
Why this matters for Ester:
Minimal test:
Data:
Metrics:
What would support the hypothesis:
What would weaken or refute the hypothesis:
Related procedures:
Related memory records:
Decision after testing:
```

Appendix B. Memory record template

```
memory_record:
  id:
  title:
  content:
  source:
  source_type: document | conversation | procedure | human_correction | judge_synthesis
  created_at:
  model_or_actor:
  procedure:
  corpus_links:
  semantic_tags:
  confidence:
  conflicts:
  status: draft | active | disputed | deprecated | superseded
  superseded_by:
  audit_log:
  encryption_scope:
  sync_scope:
```

Appendix C. Minimal control questionnaire

This questionnaire is needed to re-check the system's stability. It should be asked in the same form after 1, 7, and 30 days, preserving answers, sources, and memory statuses.

1. What is Ester in version v0.1?
2. How is Ester different from a single local LLM?
3. What does the formula `c = a + b` mean?
4. Why is the corpus limited rather than infinite?
5. What counts as a parental procedure, and what is only a metaphor?
6. How should the system handle a conflict between a new record and prior memory?
7. Which claims does v0.1 prohibit?
8. What engineering trace is left by a document-reading procedure?
9. How does Judge help, and where are its limits?
10. What must be visible in the UI so that an answer is not a black box?

License

This work is licensed under the Creative Commons Attribution 4.0 International License (CC BY 4.0). See `LICENSE` in the repository root.

DOI: 10.5281/zenodo.20679718